


Apache NiFi

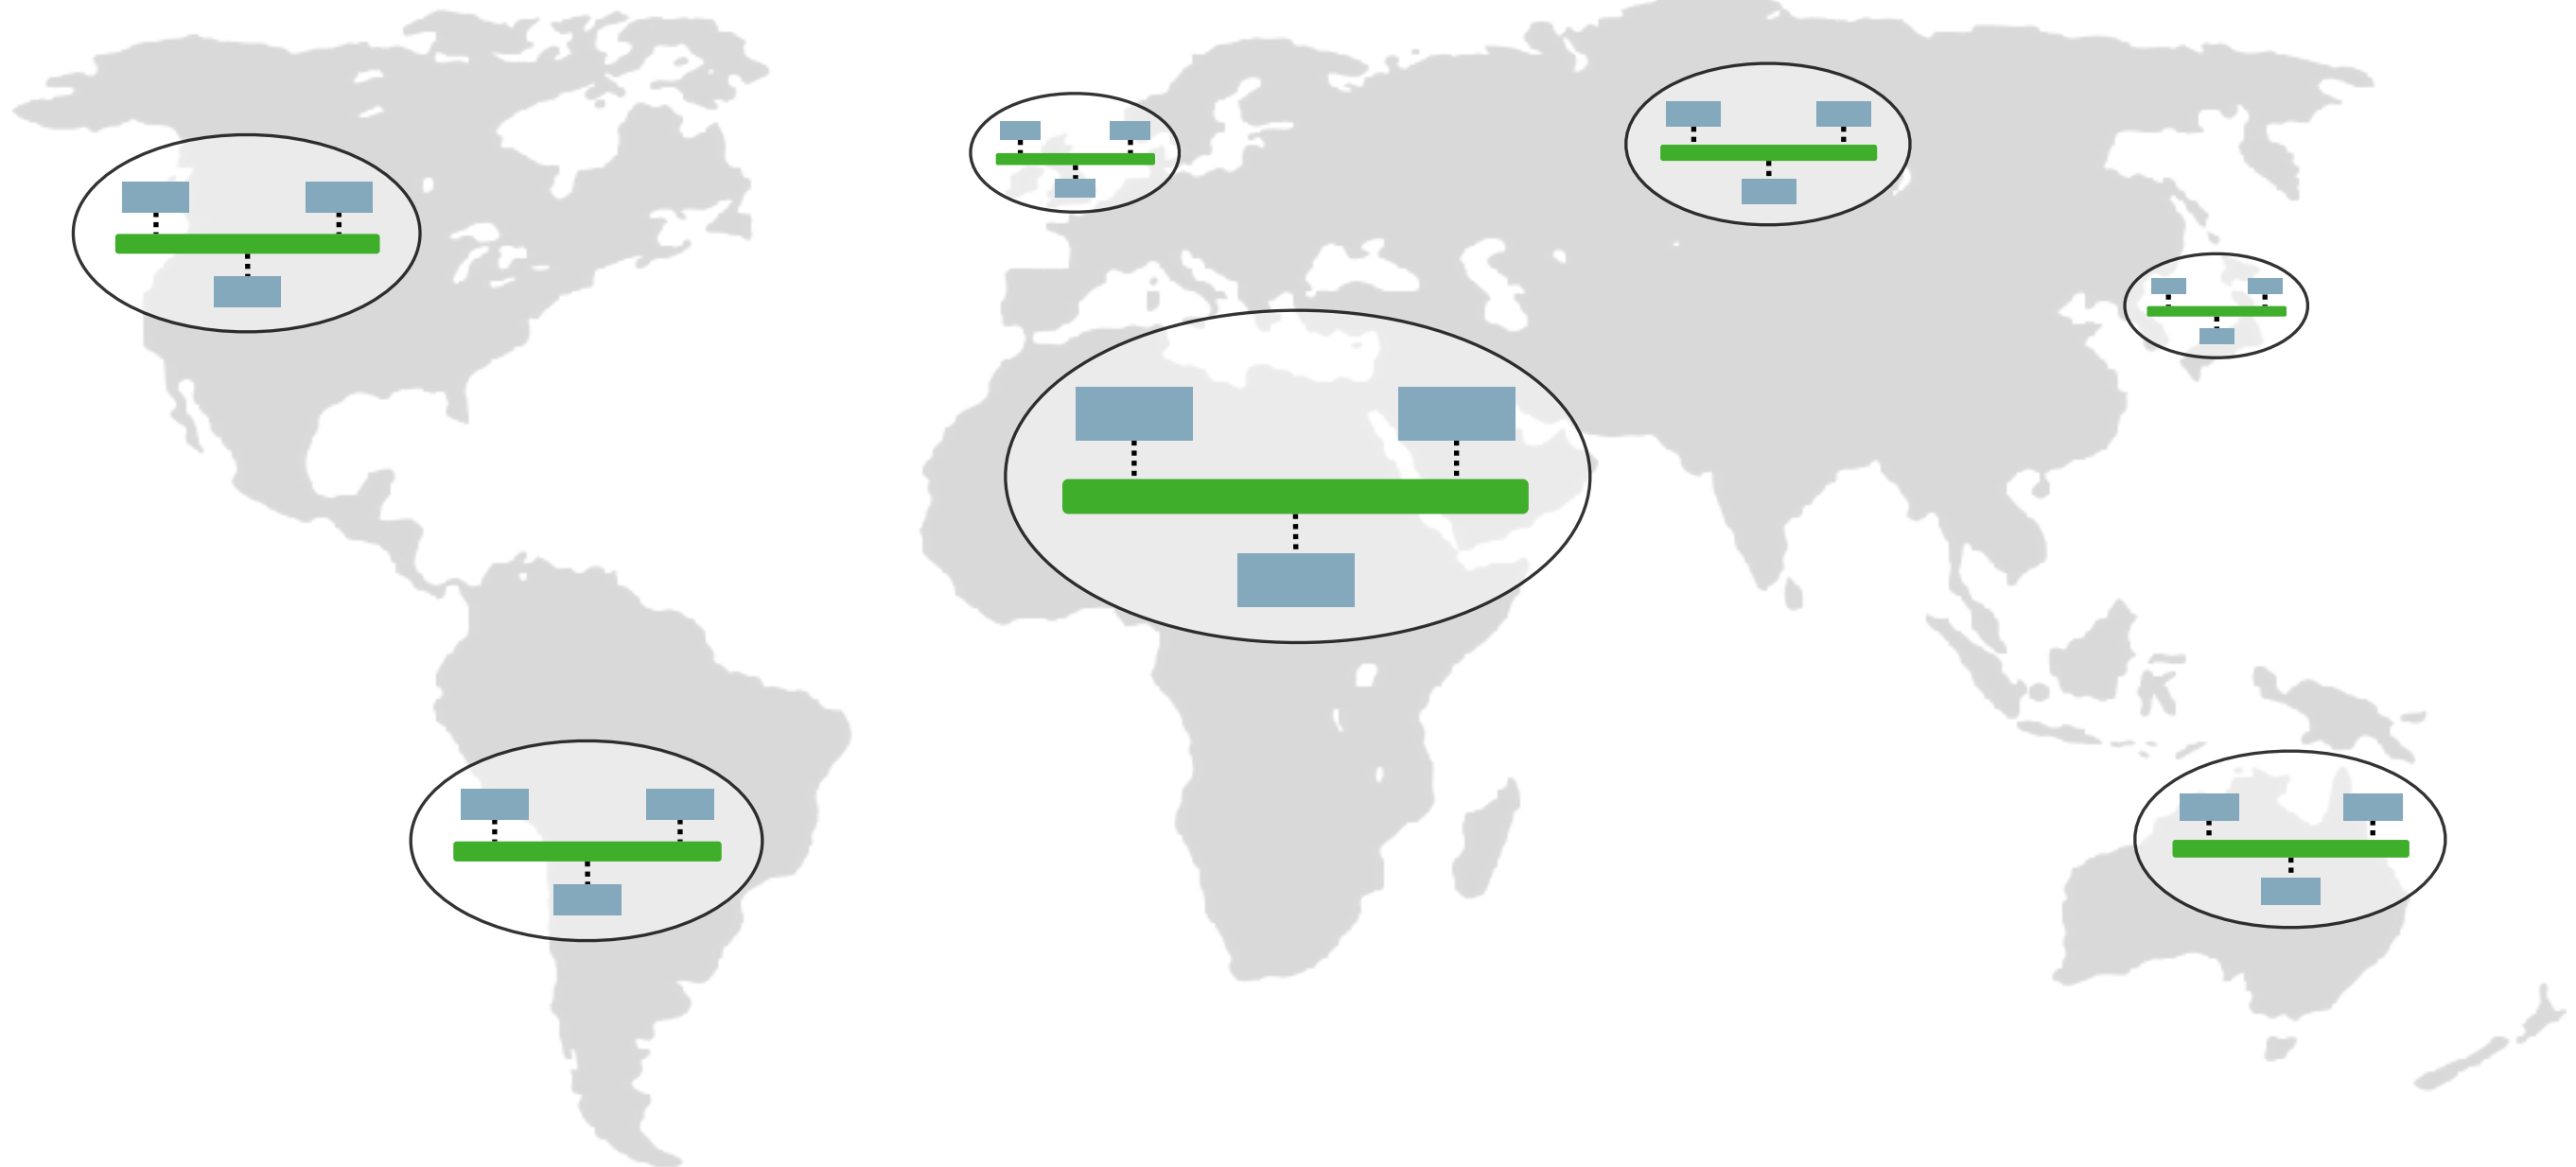
What is Apache NiFi?

Simplistic View of Enterprise Data Flow



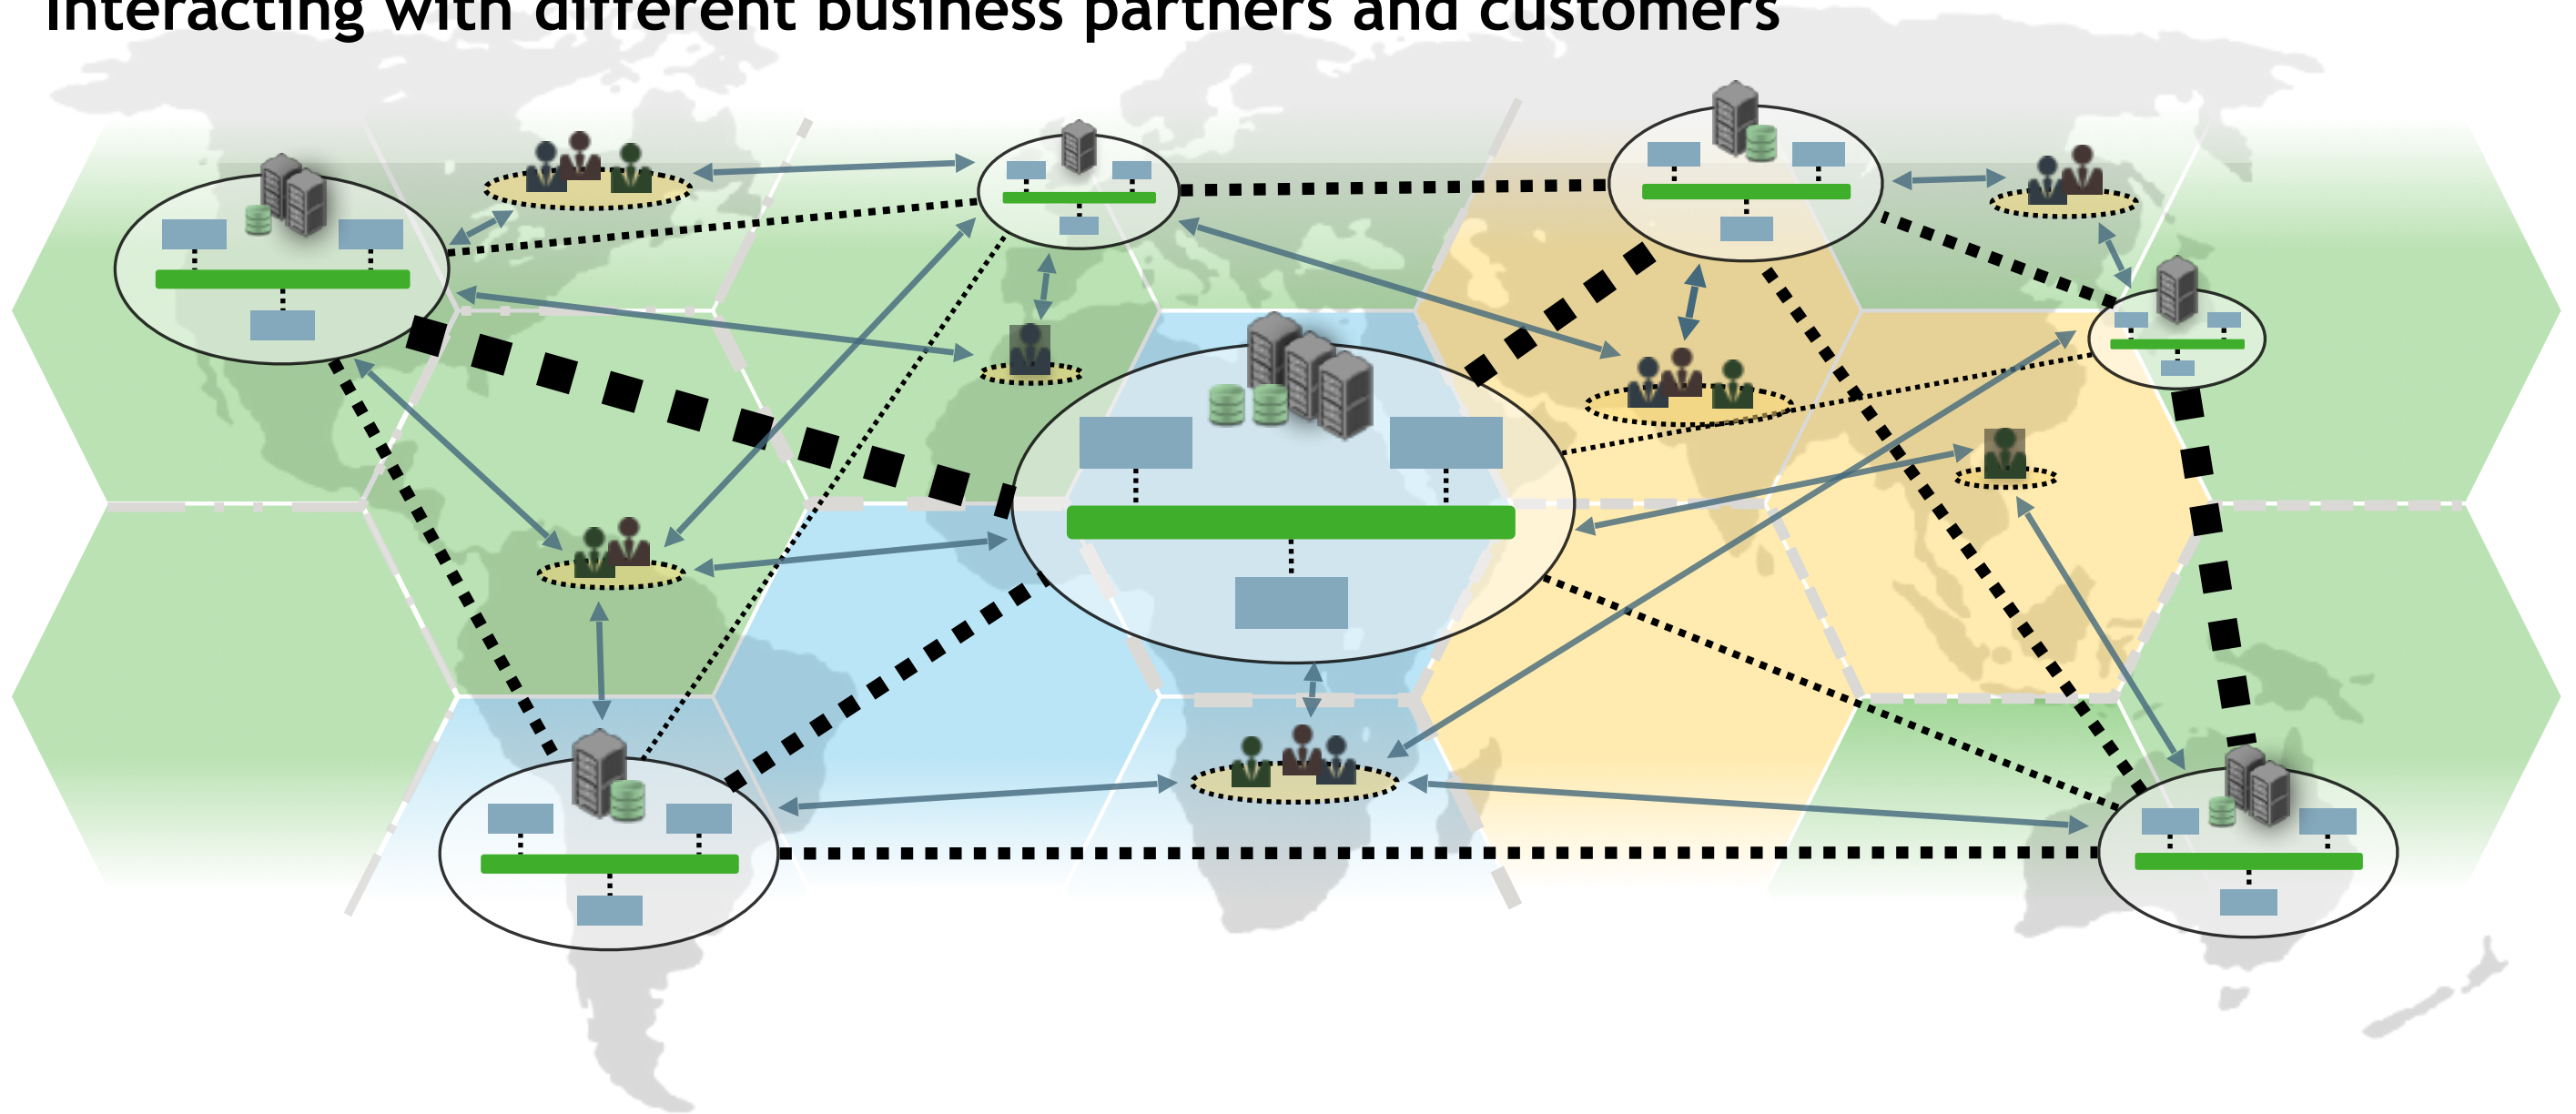
Realistic View of Enterprise Data Flow

Different organizations/business units across different geographic locations...



Realistic View of Enterprise Data Flow

Interacting with different business partners and customers



Apache NiFi

- Created to address the challenges of global enterprise dataflow
- Key features:
 - Visual Command and Control
 - Data Lineage (Provenance)
 - Data Prioritization
 - Data Buffering/Back-Pressure
 - Control Latency vs. Throughput
 - Secure Control Plane / Data Plane
 - Scale Out Clustering
 - Extensibility

Apache NiFi

What is Apache NiFi used for?

- Reliable and secure transfer of data between systems
- Delivery of data from sources to analytic platforms
- Enrichment and preparation of data:
 - Conversion between formats
 - Extraction/Parsing
 - Routing decisions

What is Apache NiFi NOT used for?

- Distributed Computation
- Complex Event Processing
- Joins / Complex Rolling Window Operations

Hadoop Ecosystem Integrations

HDFS Ingest

MergeContent

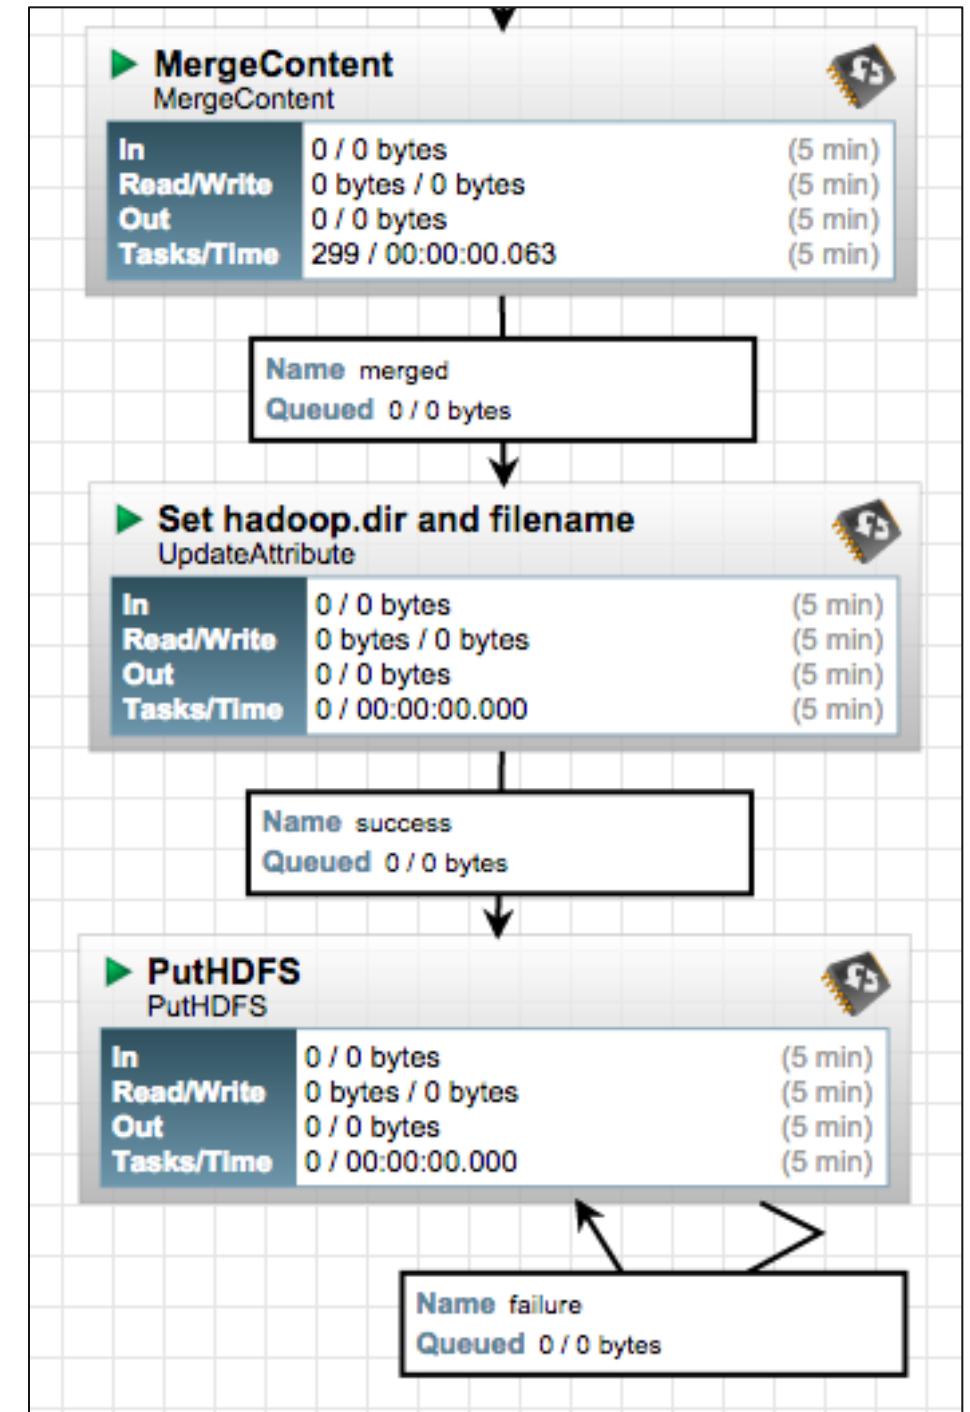
- Merges into appropriately sized files for HDFS
- Based on size, number of messages, and time

UpdateAttribute

- Sets the HDFS directory and filename
- Use expression language to dynamically bin by date:
`/data/${now():format('yyyy/MM/dd/HH')}/`

PutHDFS

- Writes FlowFile content to HDFS
- Supports Conflict Resolution Strategy and Kerberos authentication



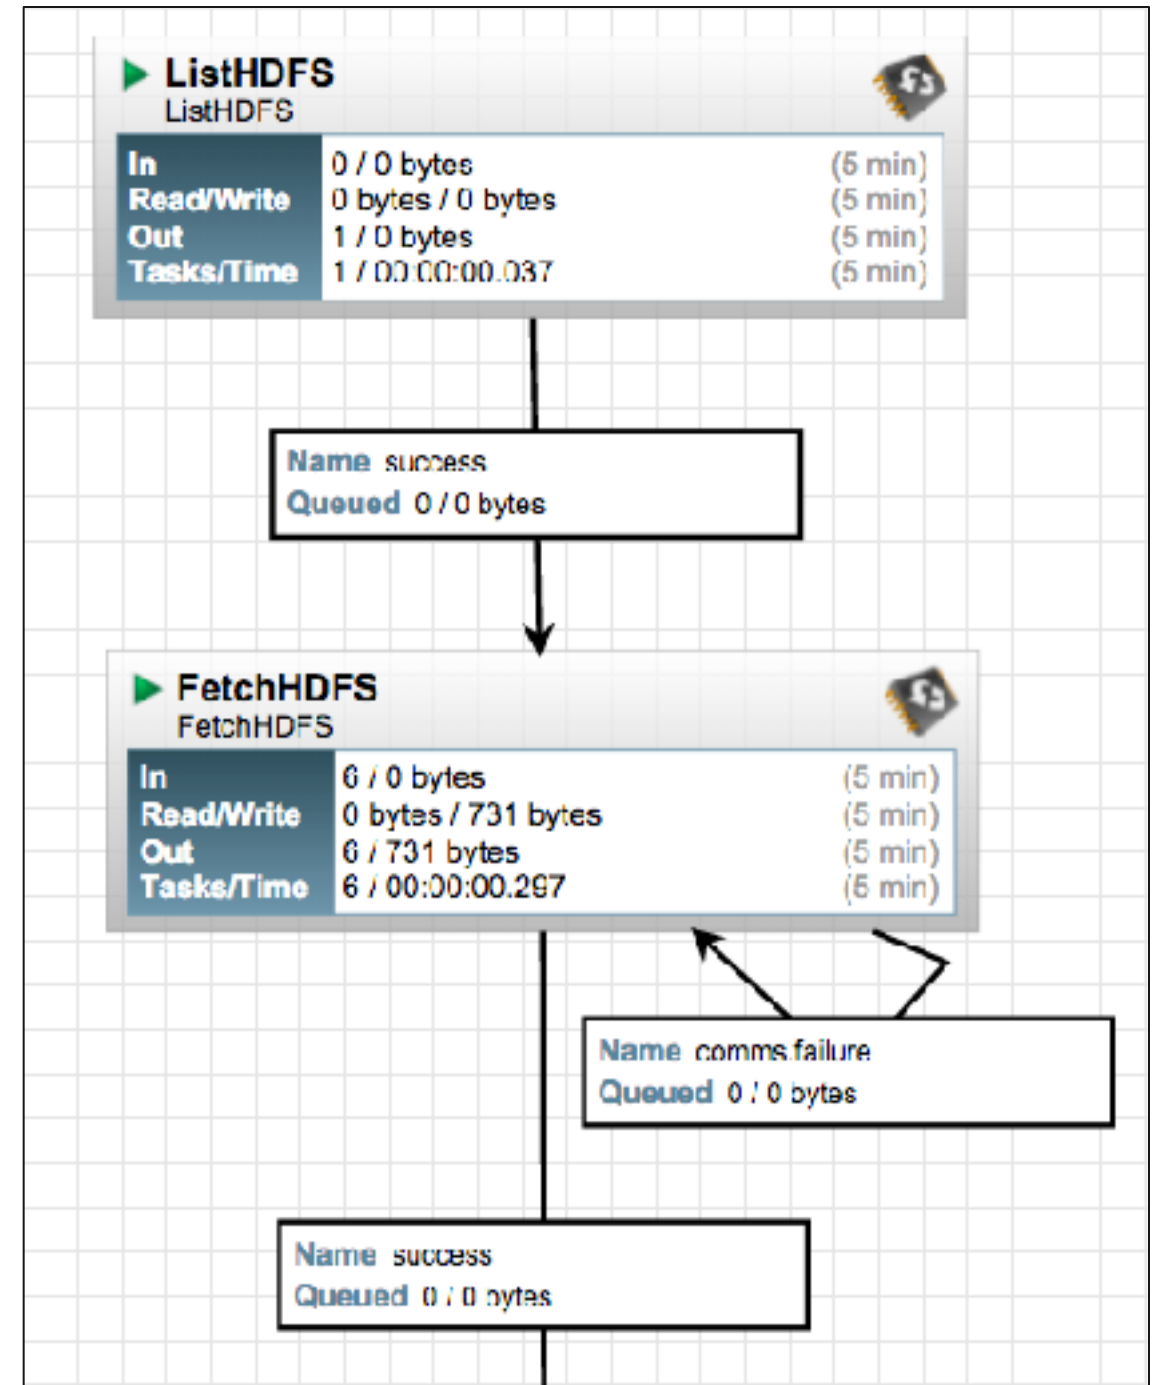
HDFS Retrieval

ListHDFS

- Periodically perform listing on HDFS directory
- Produces FlowFile per HDFS file
- Flow only contains HDFS path & filename

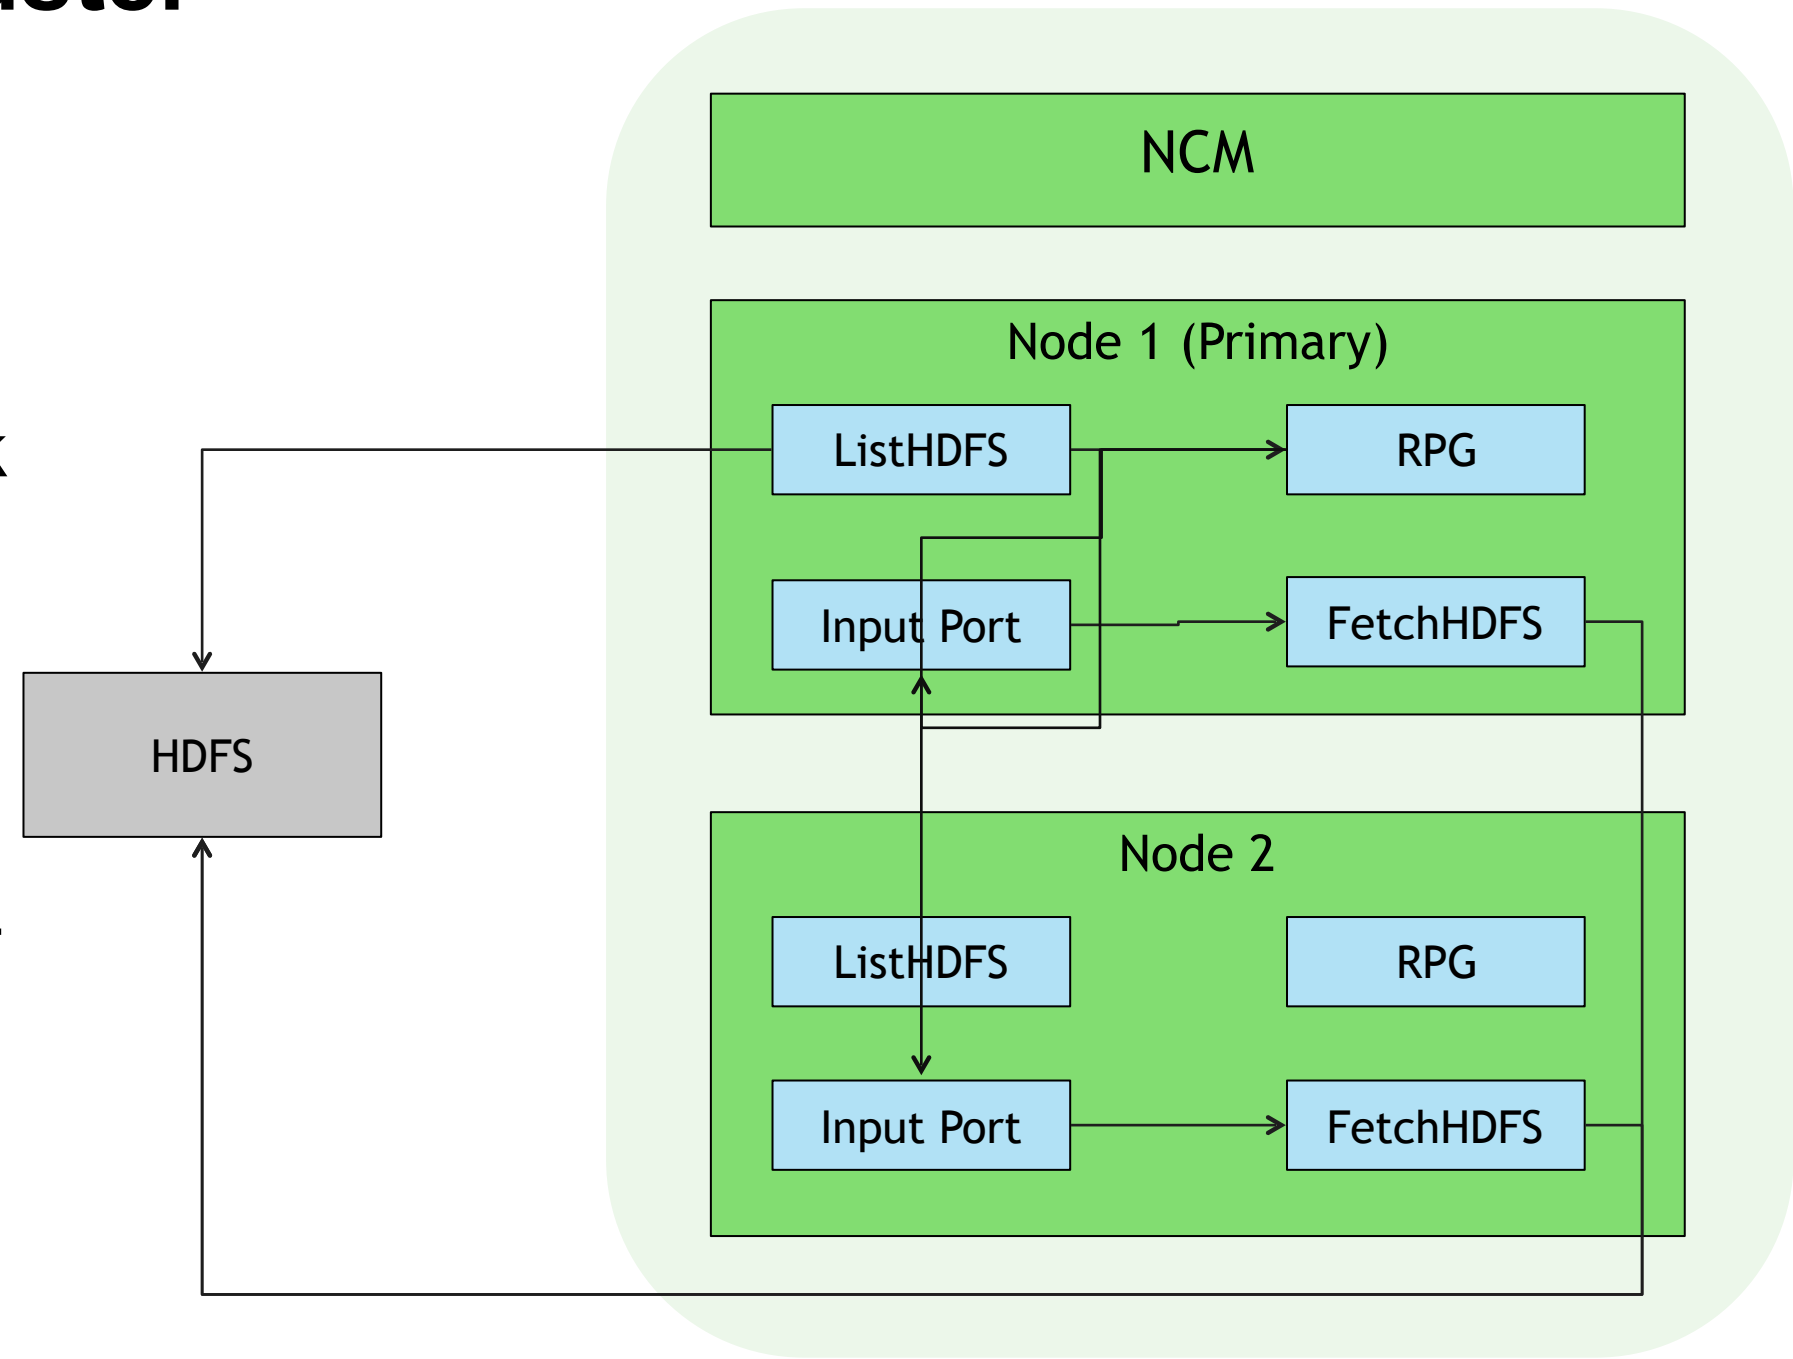
FetchHDFS

- Retrieves a file from HDFS
- Use incoming FlowFiles to dynamically fetch:
HDFS Filename: `${path}/${filename}`



HDFS Retrieval in a Cluster

- Perform “list” operation on primary node
- Send results to Remote Process Group pointing back to same cluster
- Redistributes results to all nodes to perform “fetch” in parallel
- Same approach for ListFile + FetchFile and ListSFTP + FetchSFTP



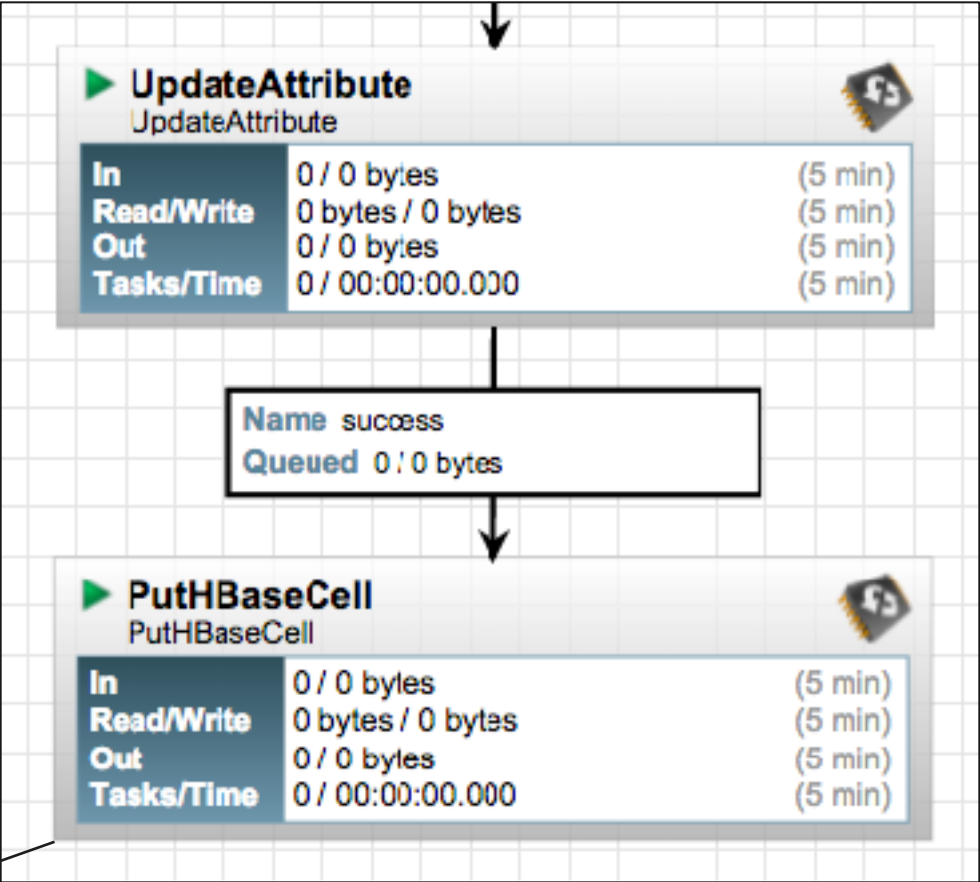
HBase Integration

- ControllerService wrapping HBase Client
- Implementation provided for HBase 1.1.2 Client
- Other implementations could be built as an extension

Property		Value
Hadoop Configuration Files	?	/etc/hbase/conf/hbase-site.xml,/etc/hadoop/conf/core-si...
Kerberos Principal	?	No value set
Kerberos Keytab	?	No value set
ZooKeeper Quorum	?	No value set
ZooKeeper Client Port	?	No value set
ZooKeeper ZNode Parent	?	No value set
HBase Client Retries	?	1

HBase Ingest – Single Cell

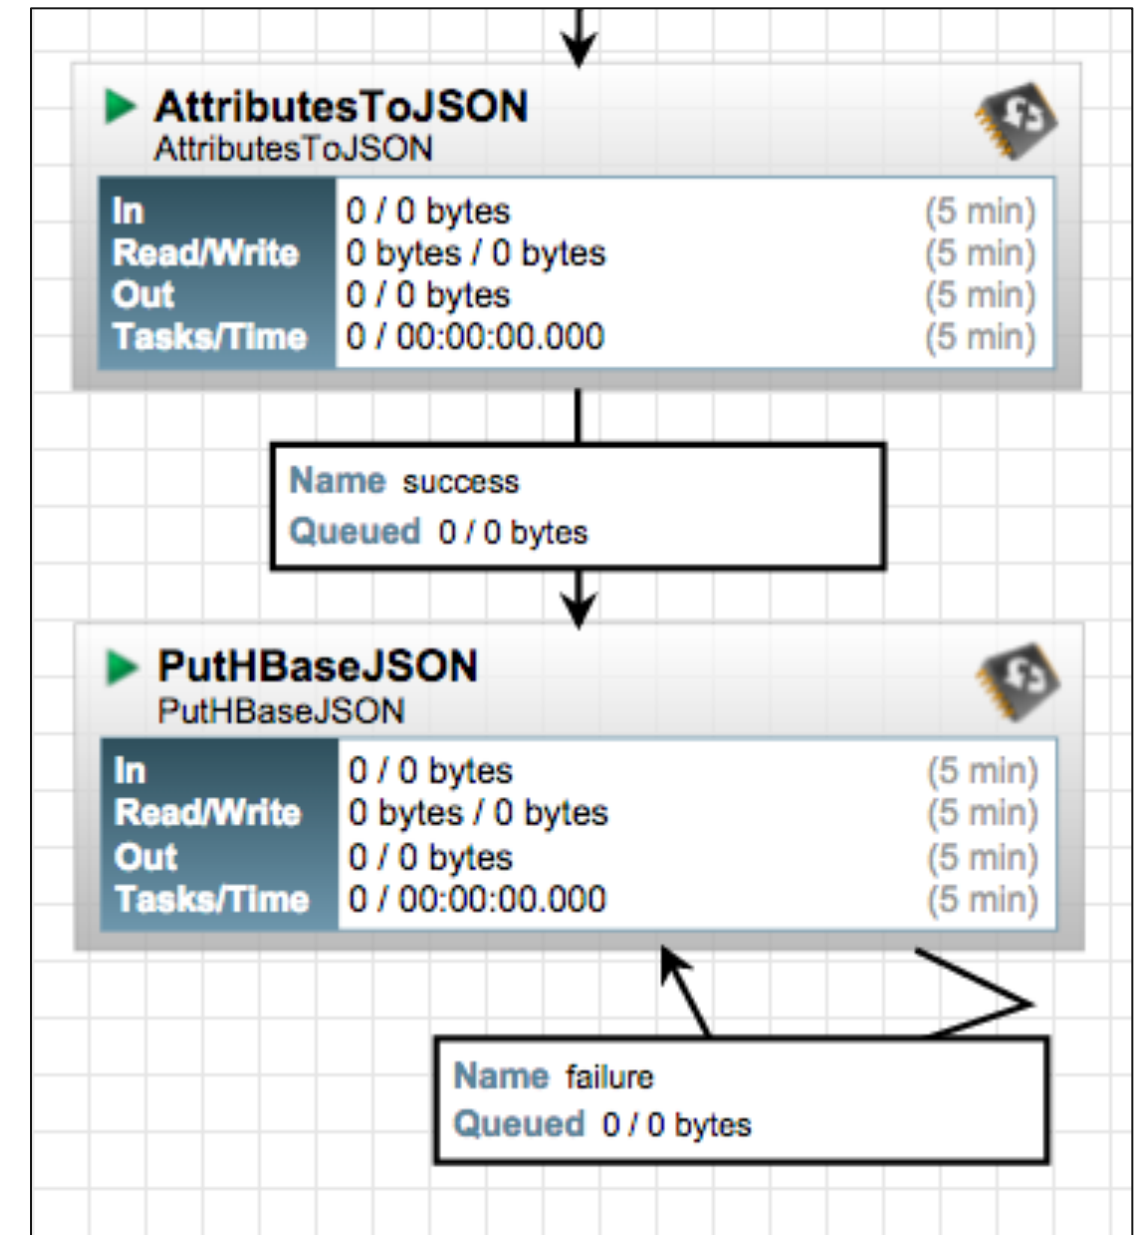
- Table, Row Id, Col Family, and Col Qualifier provided in processor, or dynamically from attributes
- FlowFile content becomes the cell value
- Batch Size to specify maximum number of cells for a single 'put' operation



Property		Value
HBase Client Service	?	HBase_1_1_2_ClientService
Table Name	?	\${hbase.table}
Row Identifier	?	\${hbase.row}
Column Family	?	\${hbase.cf}
Column Qualifier	?	\${hbase.cq}
Batch Size	?	25

HBase Ingest – Full Row

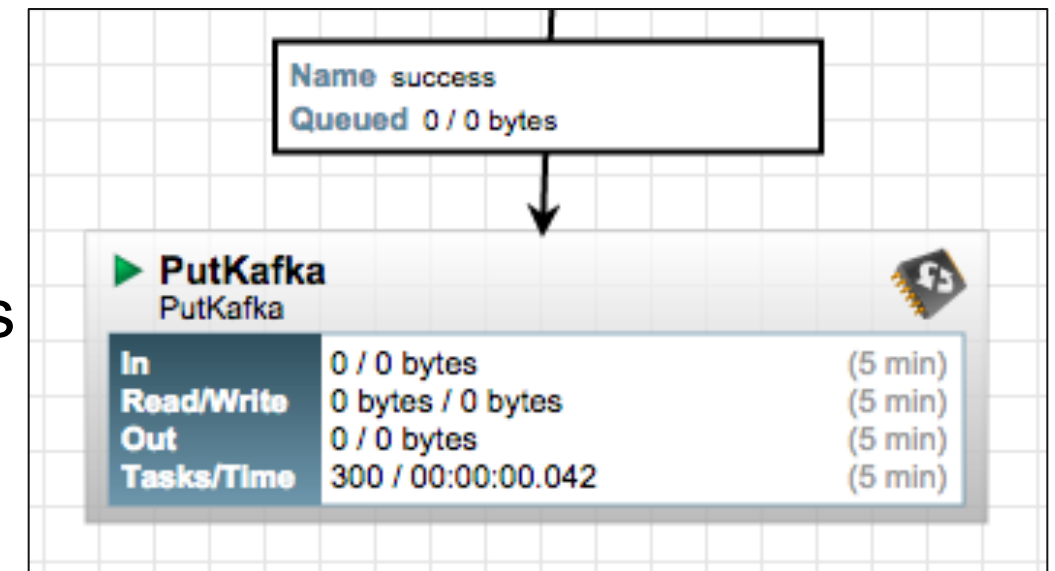
- Table and Column Family provided in processor, or dynamically from attributes
- Row ID can be a field in JSON, or a FlowFile attribute
- JSON Field/Values become Column Qualifiers and Values
- Complex Field Strategy
 - Fail
 - Warn
 - Ignore
 - Text



Kafka Integration

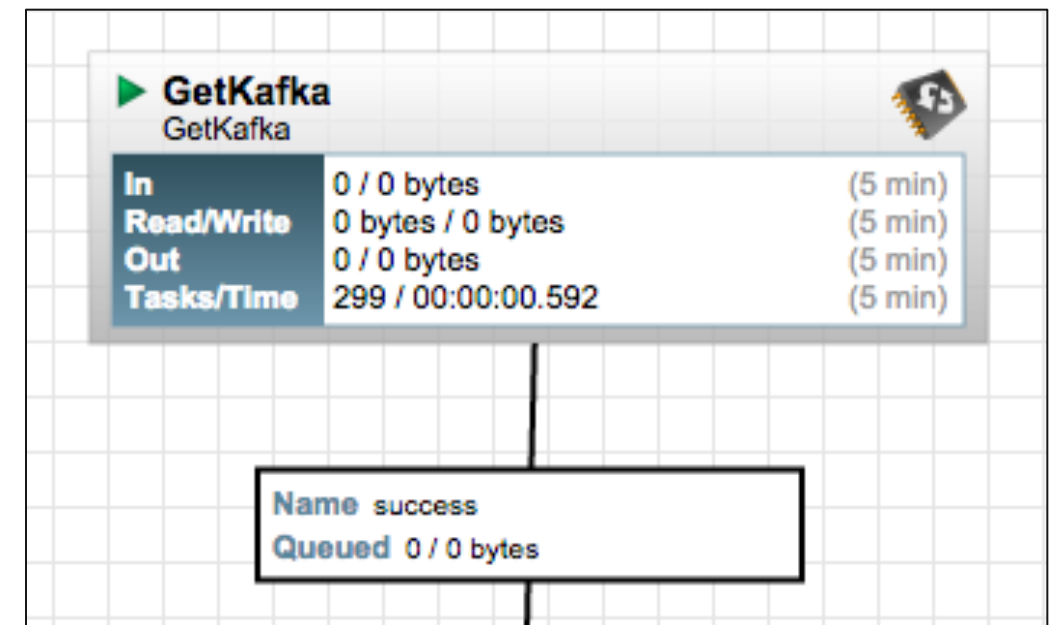
PutKafka

- Provide Broker and Topic Name
- Publishes FlowFile content as one or more messages
- Ability to send large delimited content, split into messages by NiFi



GetKafka

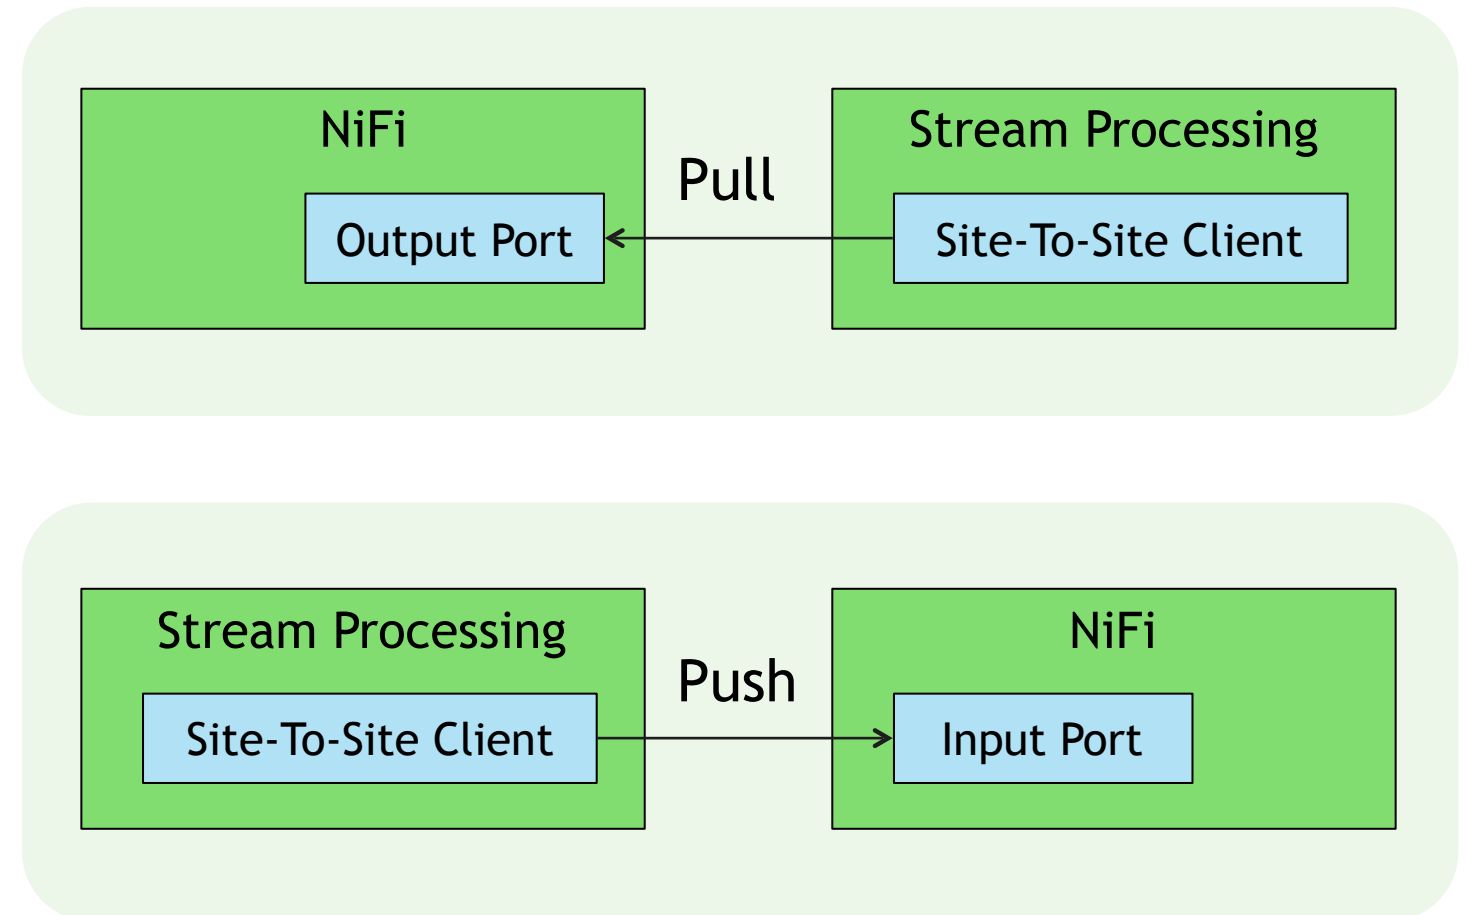
- Provide ZK Connection String and Topic Name
- Produces a FlowFile for each message consumed



Stream Processing Integration

- Stream processing systems generally pull data, then push results
- NiFi Site-To-Site pushes and pulls between NiFi instances
- The Site-To-Site Client can be used from a stream processing platform

<https://github.com/apache/nifi/tree/master/nifi-commons/nifi-site-to-site-client>



Site-to-Site Client Overview

- Push to Input Port, or Pull from Output Port
- Communicate with NiFi clusters, or standalone instances
- Handles load balancing and reliable delivery
- Secure connections using certificates (optional)

```
SiteToSiteClientConfig clientConfig =  
    new SiteToSiteClient.Builder()  
        .url("http://localhost:8080/nifi")  
        .portName("My Port")  
        .buildConfig();
```

Site-To-Site Client Pulling

```
SiteToSiteClient client = ...
```

```
Transaction transaction =  
client.createTransaction(TransferDirection.RECEIVE);
```

```
DataPacket dataPacket = transaction.receive();
```

```
while (dataPacket != null) {
```

```
...
```

```
}
```

```
transaction.confirm();
```

```
transaction.complete();
```

Site-To-Site Client Pushing

```
SiteToSiteClient client = ...
```

```
Transaction transaction =  
client.createTransaction(TransferDirection.SEND);
```

```
NiFiDataPacket data = ...
```

```
transaction.send(data.getContent(), data.getAttributes());  
transaction.confirm();
```

```
transaction.complete();
```

Current Stream Processing Integrations

Spark Streaming - NiFi Spark Receiver

- <https://github.com/apache/nifi/tree/master/nifi-external/nifi-spark-receiver>

Storm – NiFi Spout

- <https://github.com/apache/nifi/tree/master/nifi-external/nifi-storm-spout>

Flink – NiFi Source & Sink

- <https://github.com/apache/flink/tree/master/flink-streaming-connectors/flink-connector-nifi>

Apex - NiFi Input Operators & Output Operators

- <https://github.com/apache/incubator-apex-malhar/tree/master/contrib/src/main/java/com/datatorrent/contrib/nifi>

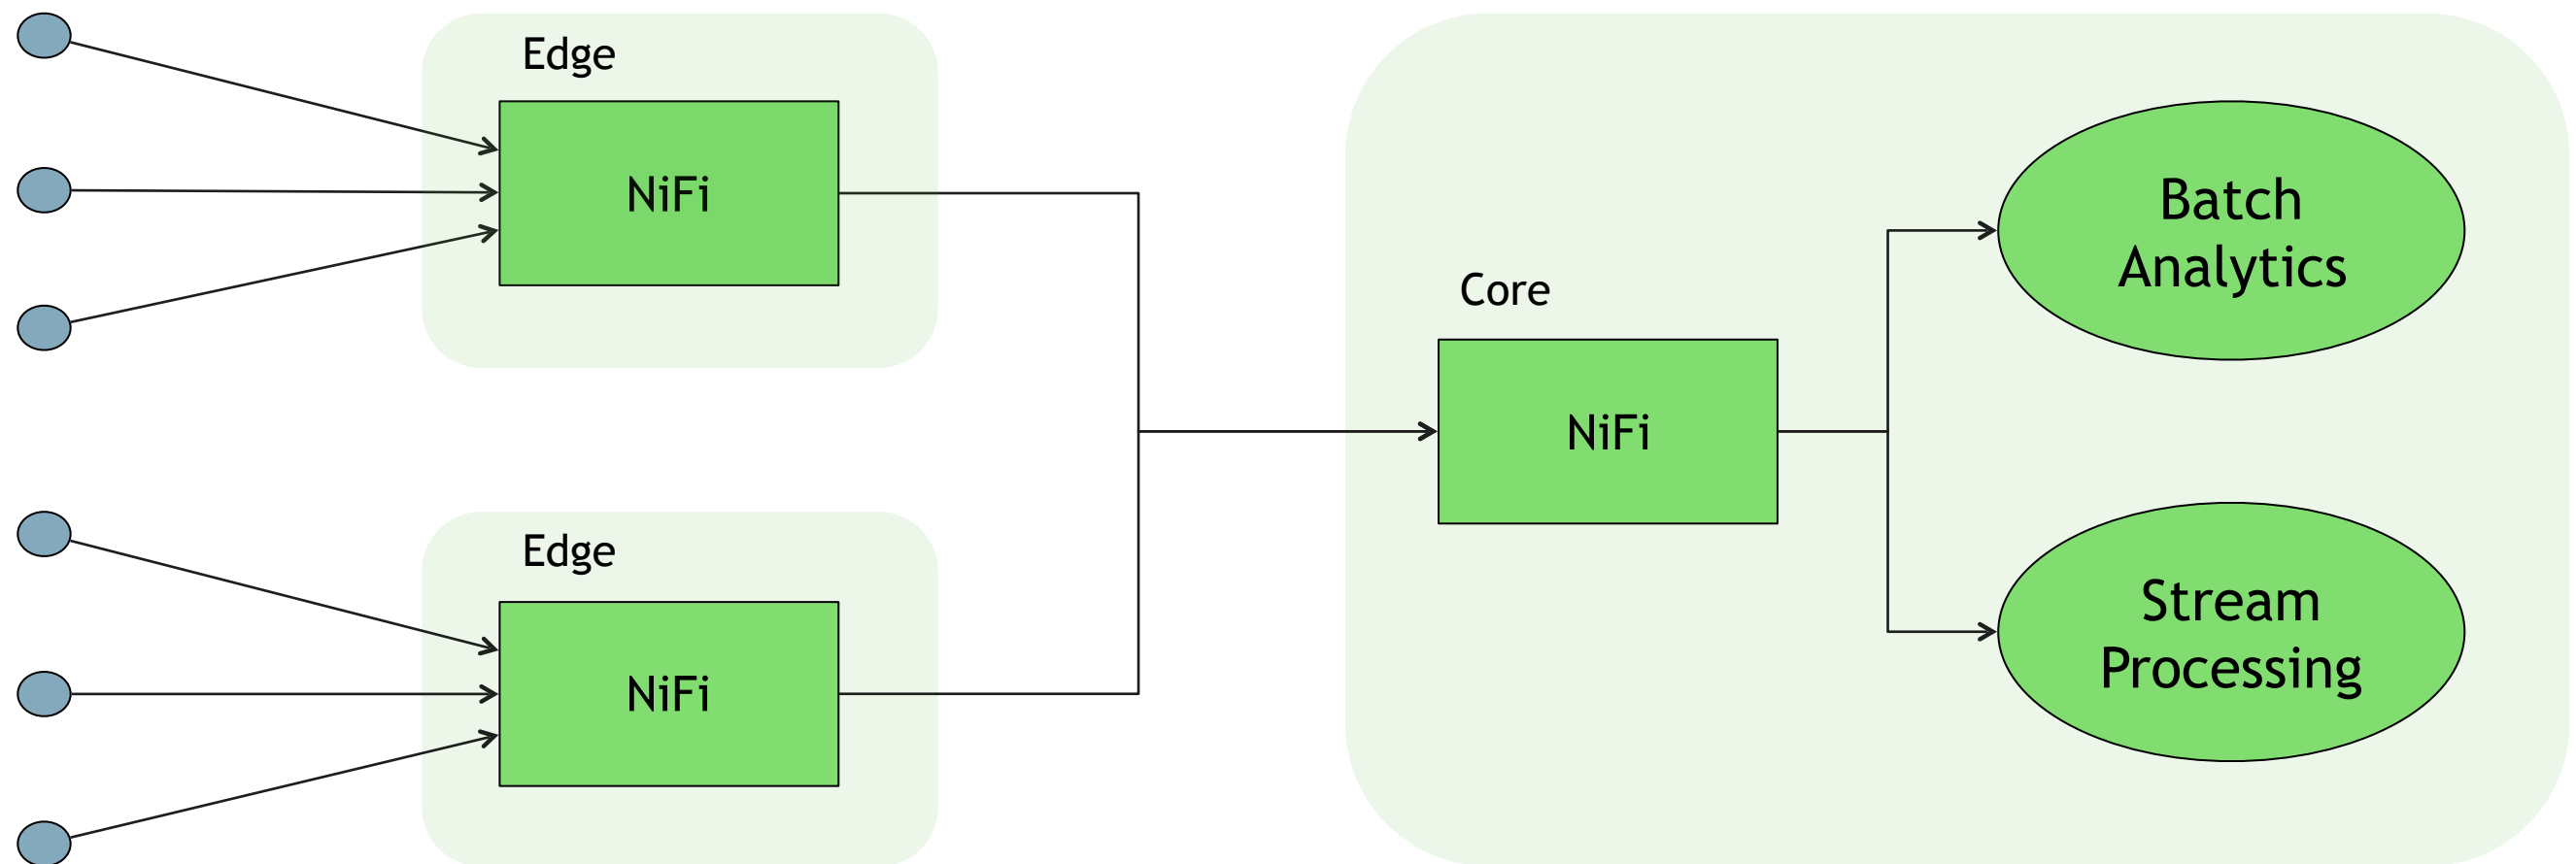
Other Relevant Integrations

- GetSolr, PutSolrContentStream
- FetchElasticSearch, PutElasticSearch
- GetMongo, PutMongo
- QueryCassandra, PutCassandraQL
- GetCouchbaseKey, PutCouchbaseKey
- QueryDatabaseTable, ExecuteSQL, PutSQL
- GetSplunk, PutSplunk
- And more! <https://nifi.apache.org/docs.html>

Use-Case/Architecture Discussion

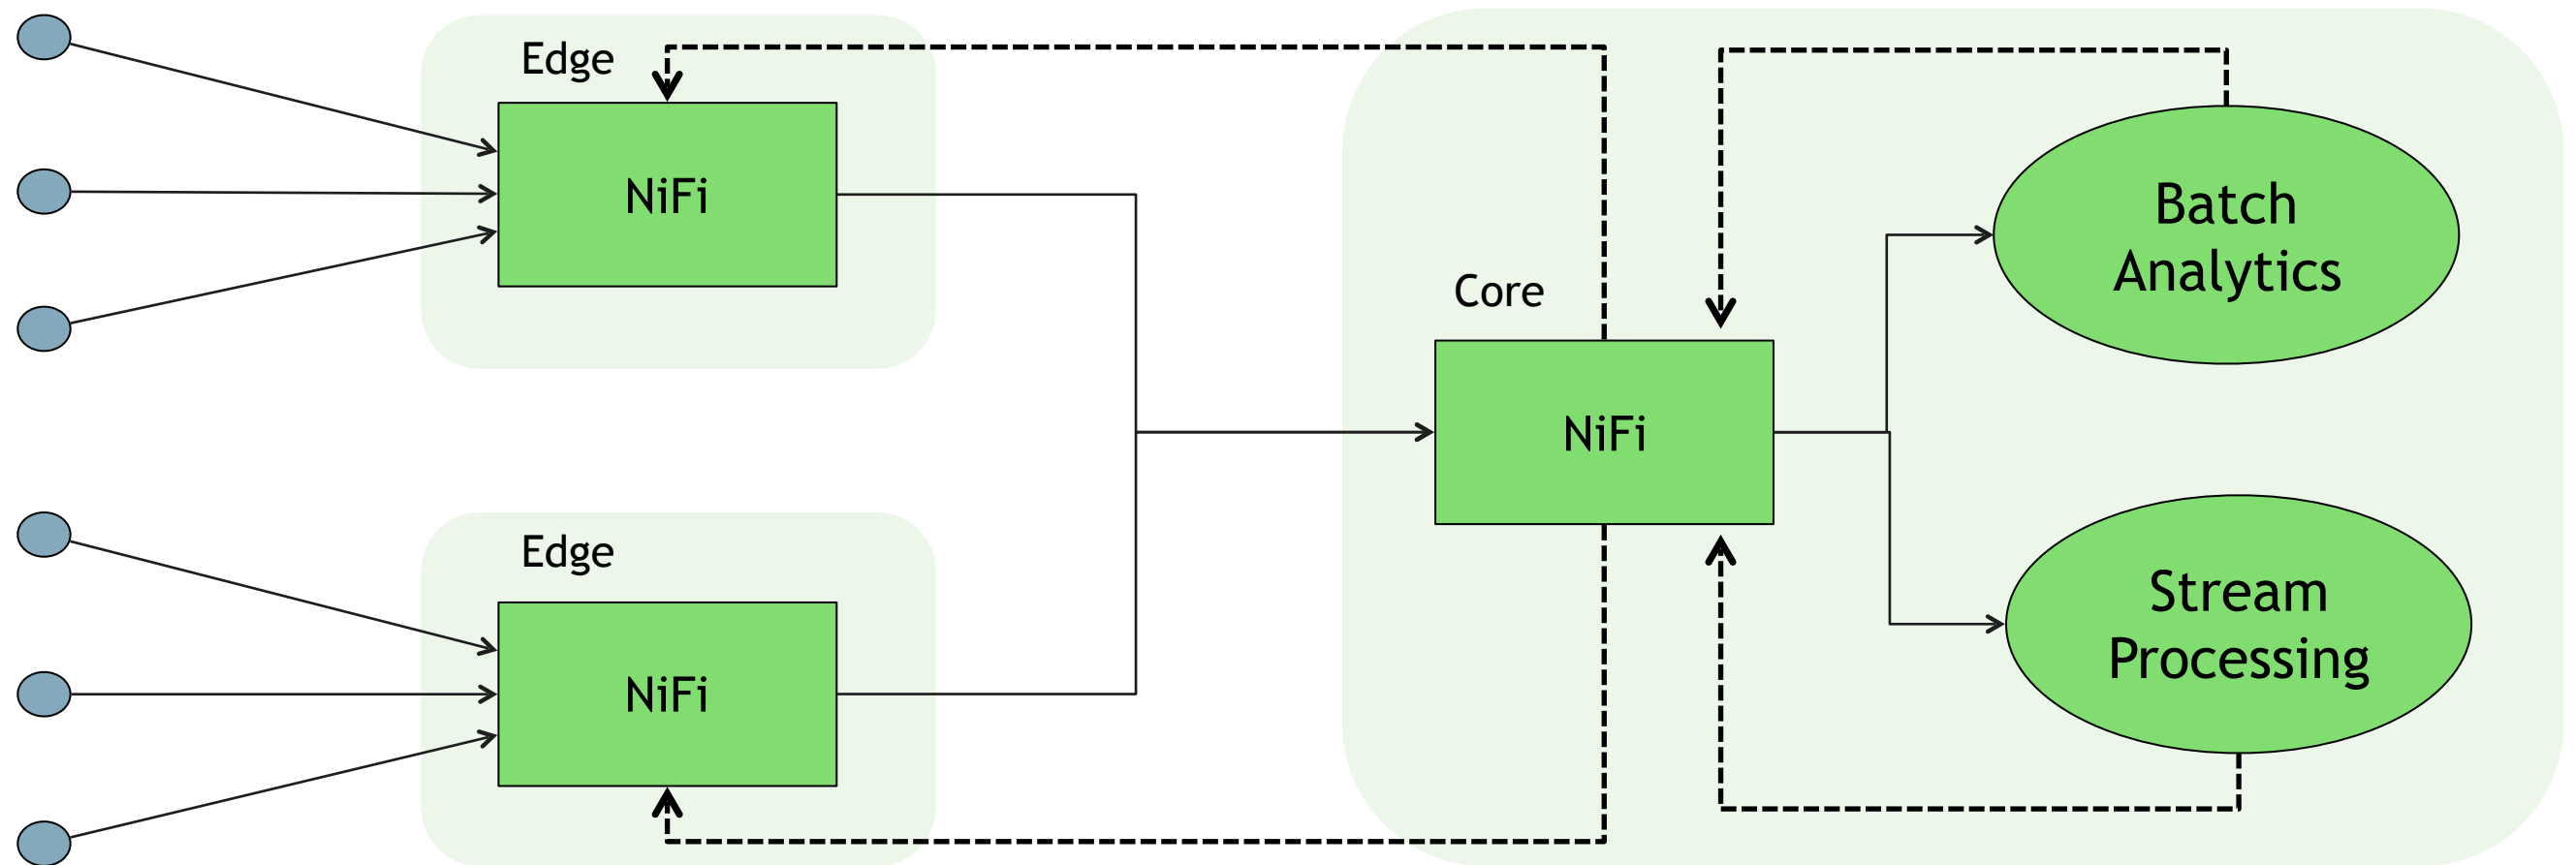
Drive Data to Core for Analysis

- Drive data from sources to central data center for analysis
- Tiered collection approach at various locations, think regional data centers



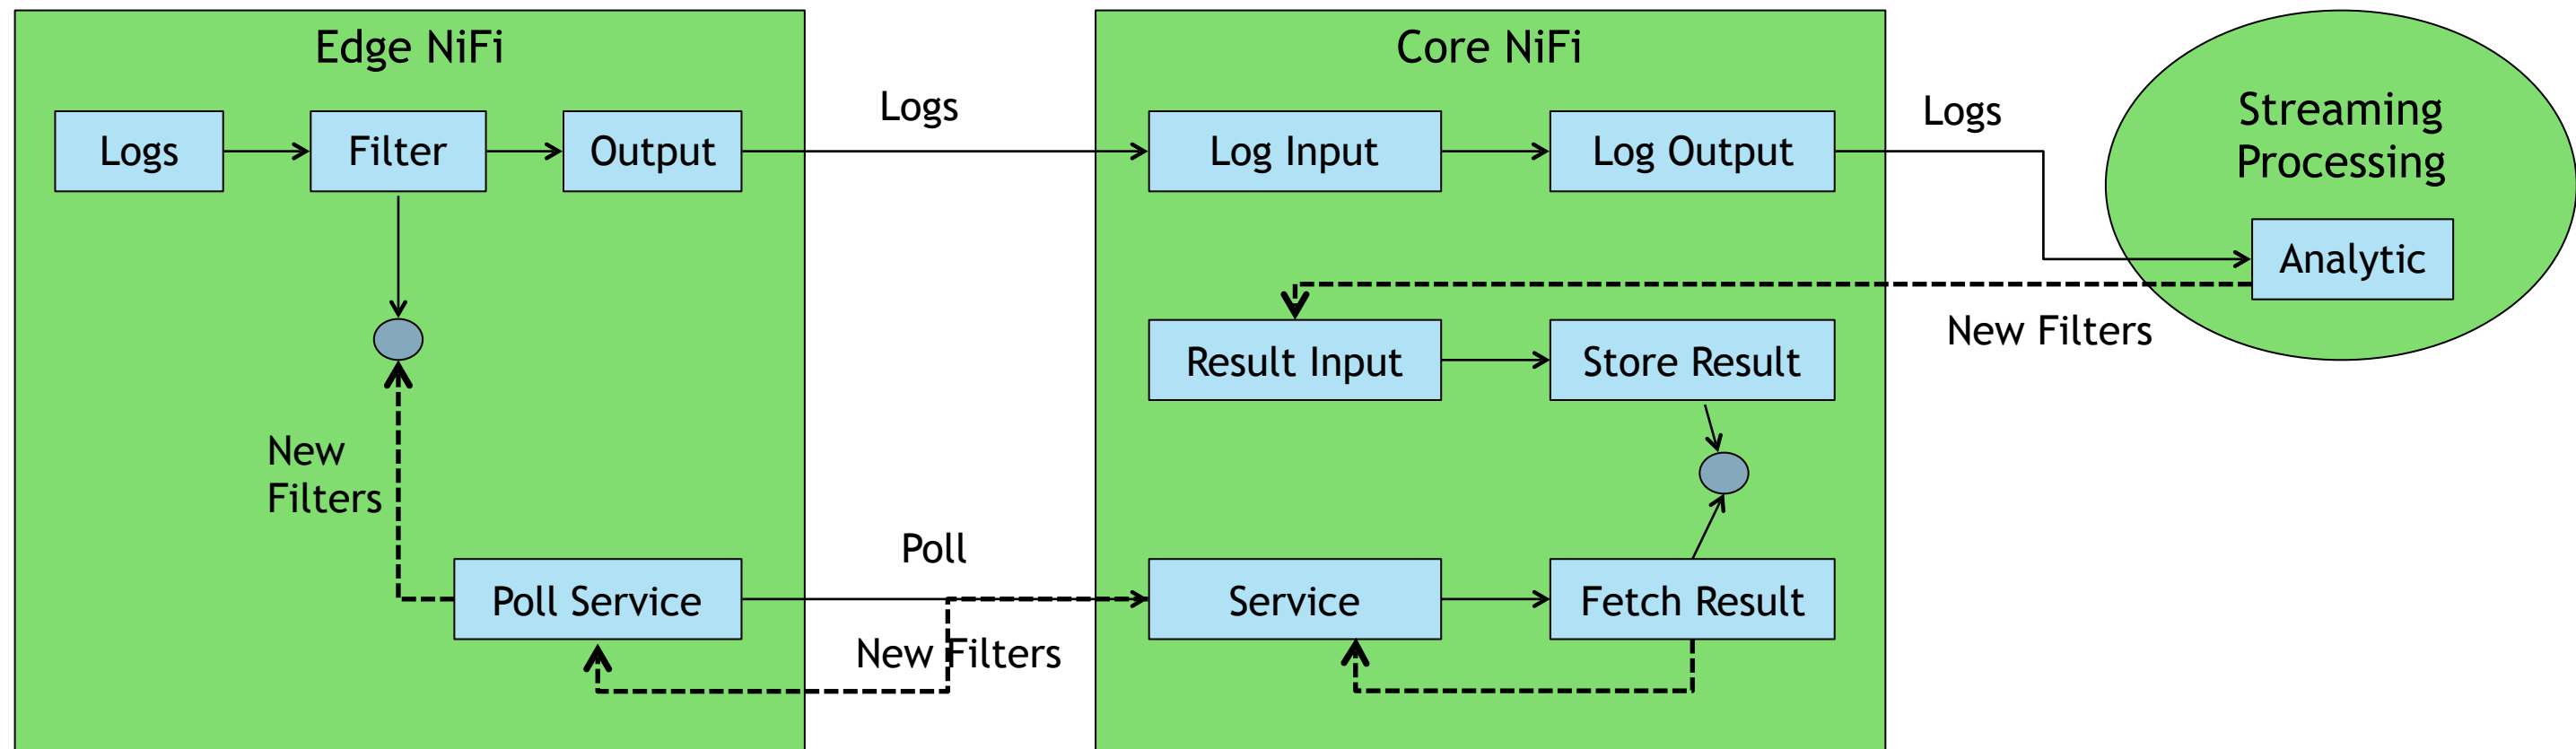
Dynamically Adjusting Data Flows

- Push analytic results back to core NiFi
- Push results back to edge locations/devices to change behavior



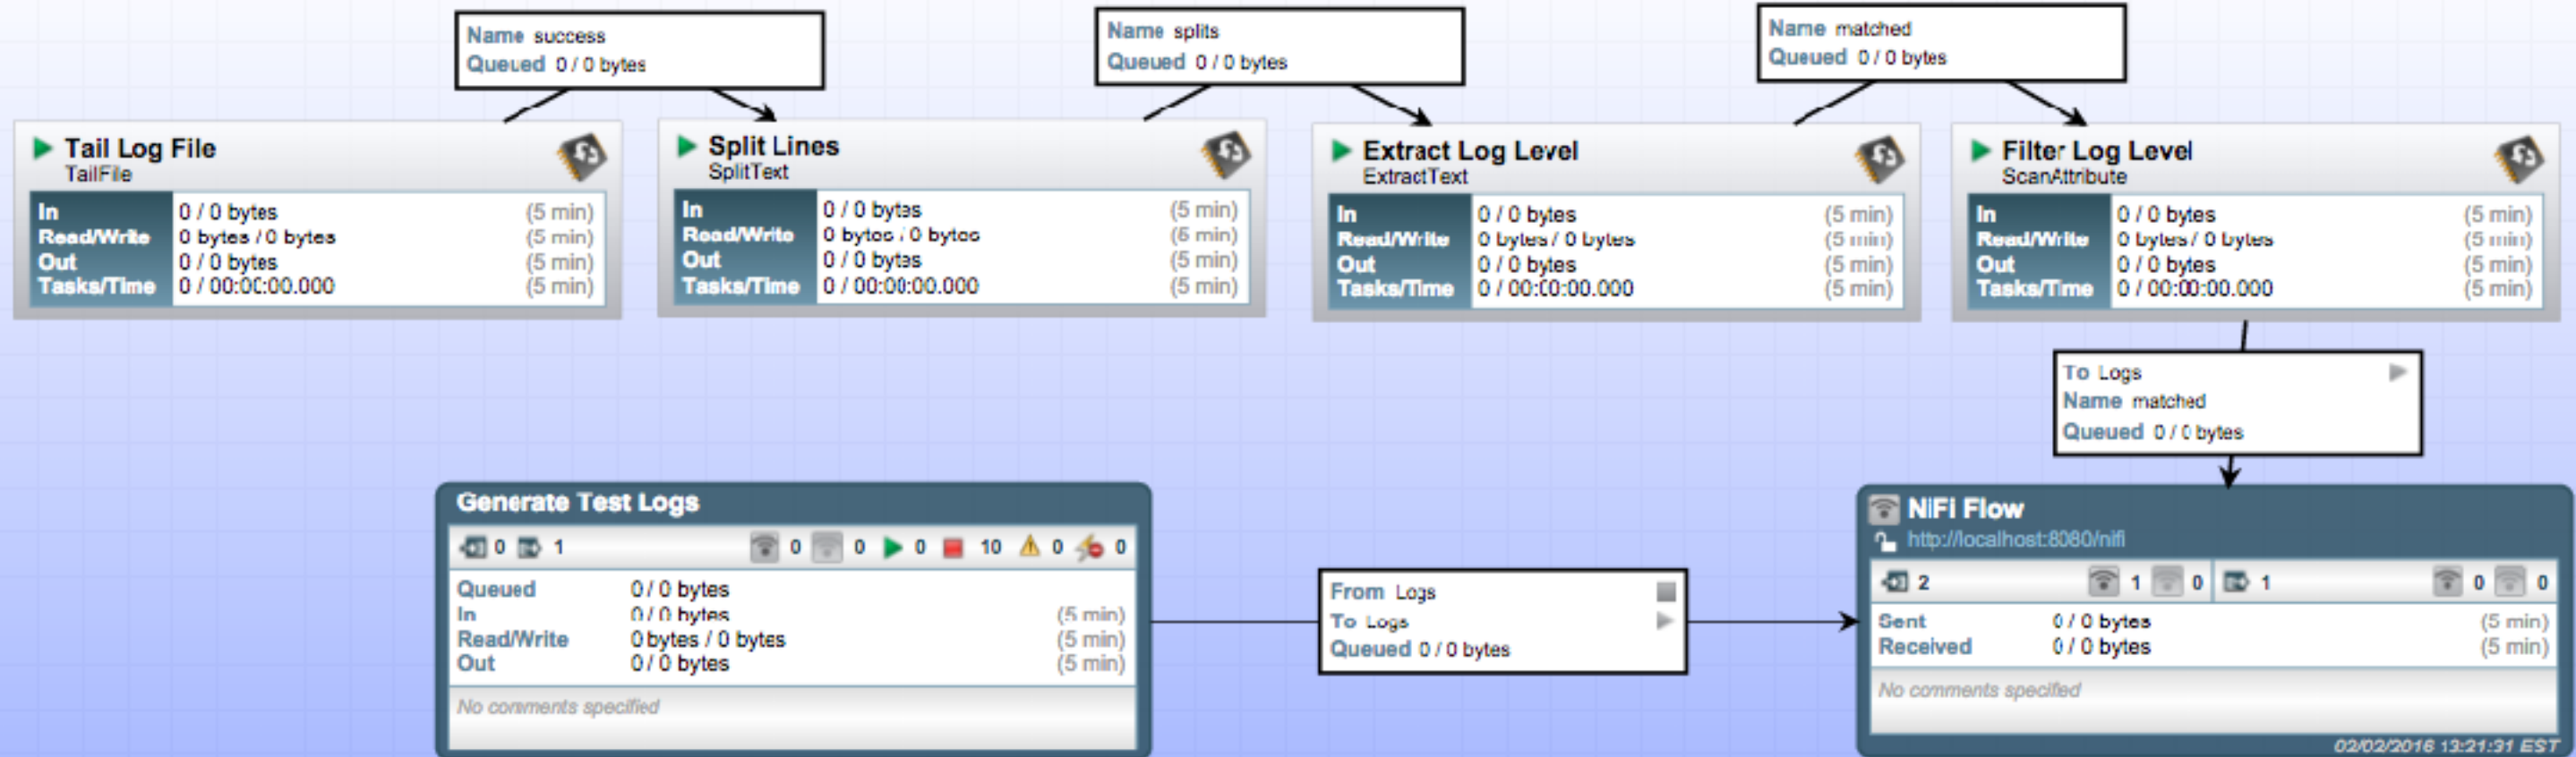
Example: Dynamic Log Collection

1. Logs filtered by level and sent from Edge -> Core
2. Stream processing produces new filters based on rate & sends back to core
3. Edge polls core for new filter levels & updates filtering




Dynamic Log Collection – Edge NiFi

Send Logs to Core




Pull New Log Level Dictionary

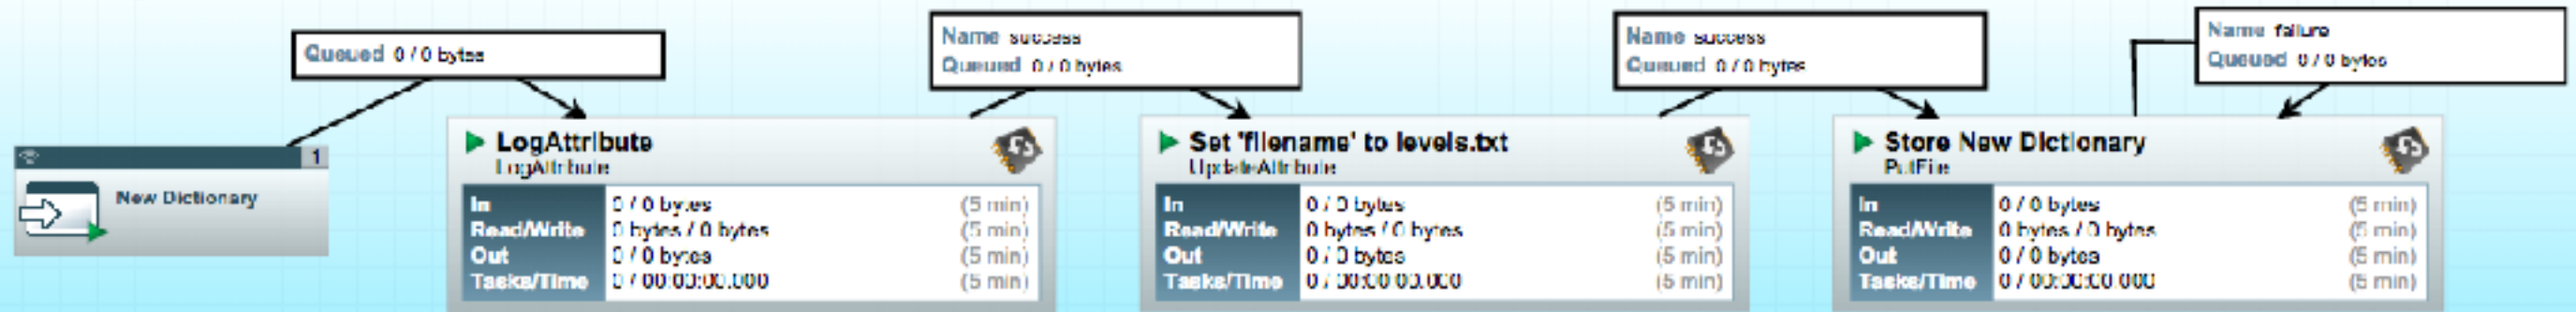


Dynamic Log Collection – Core NiFi

Receive Logs from Edge



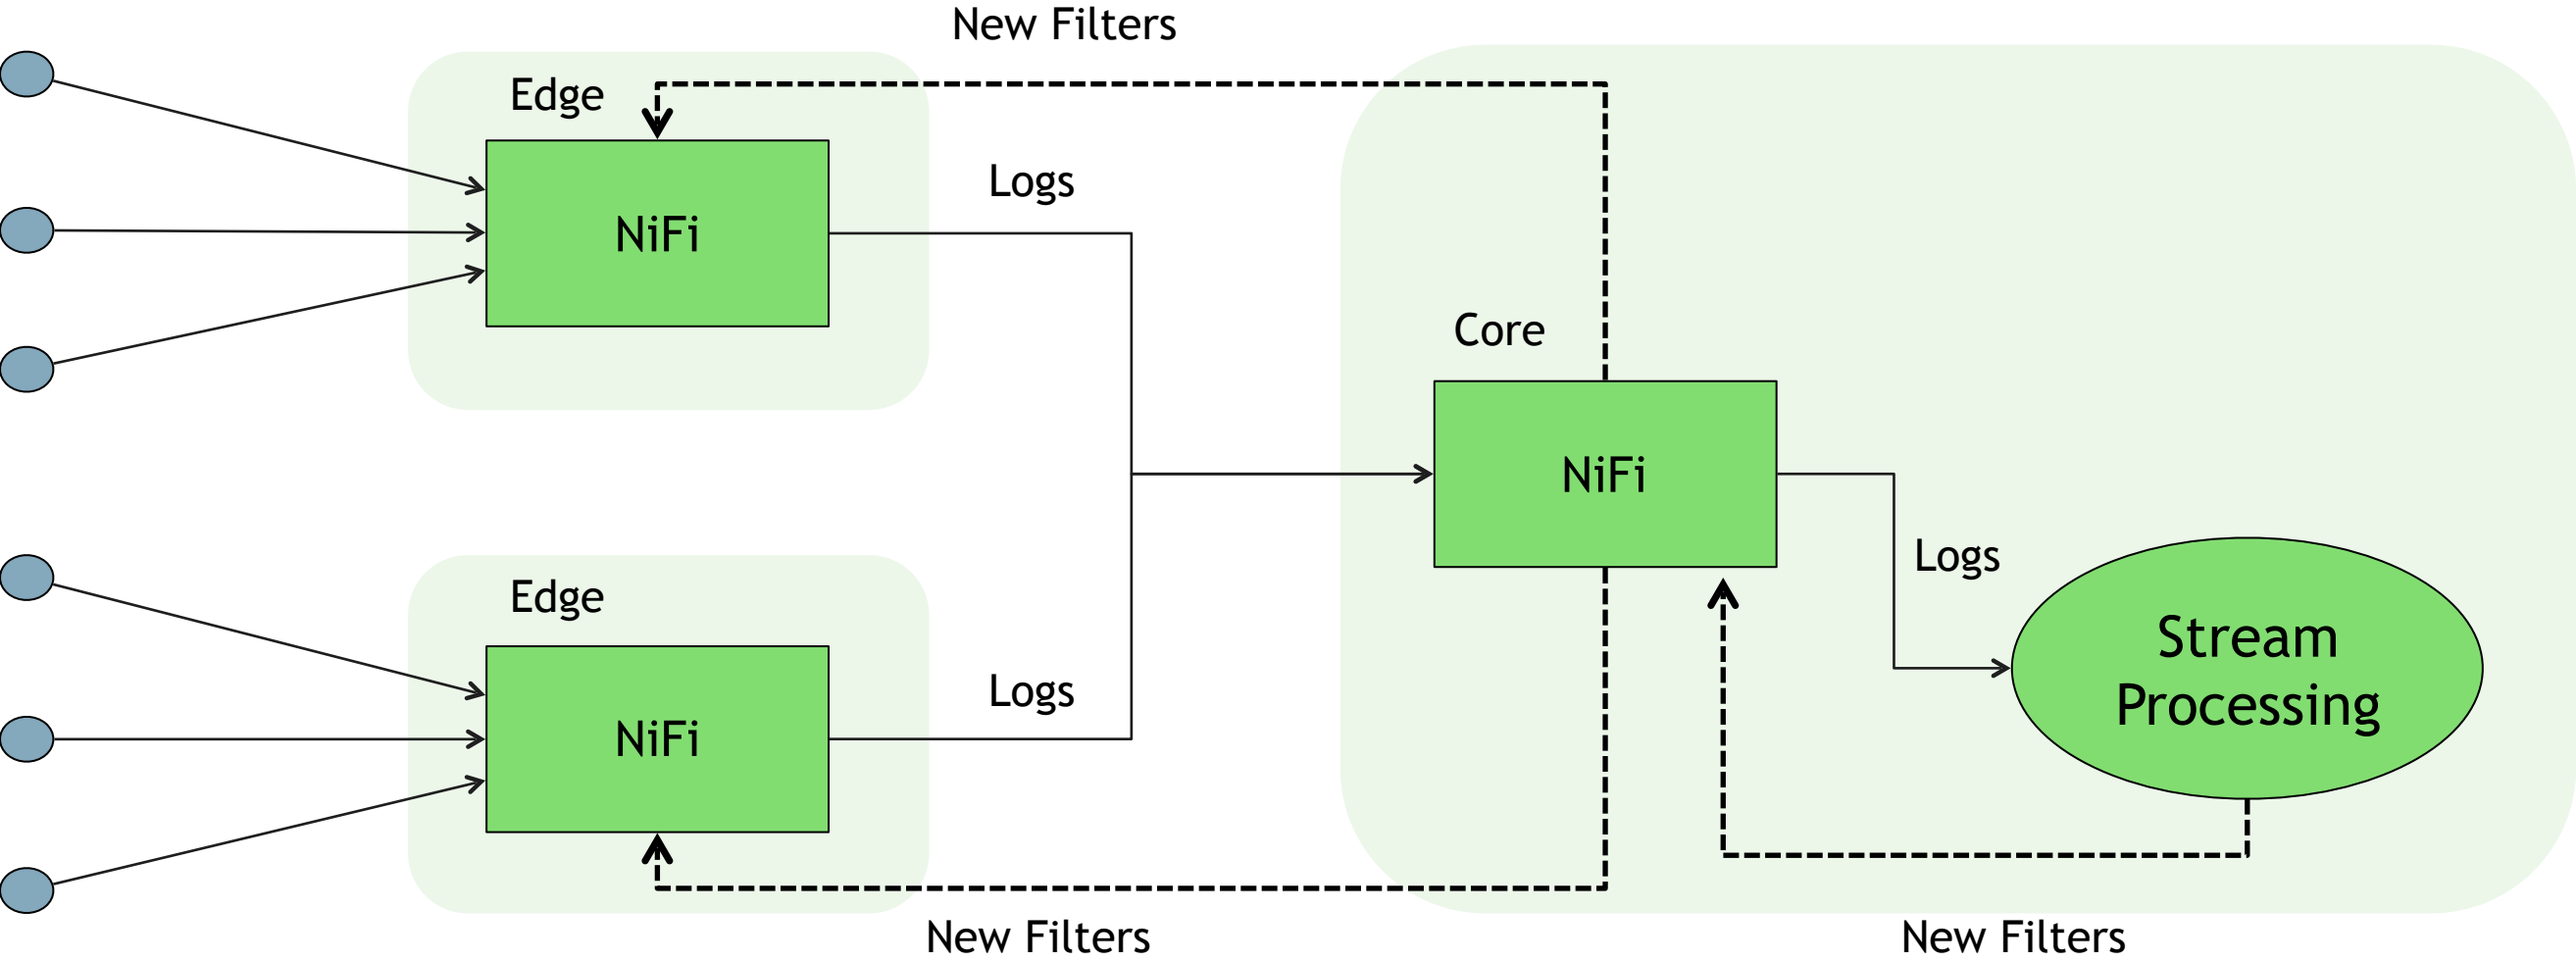
Receive Analytic Results



Handle Request for Log Levels



Dynamic Log Collection Summary



The Future – Ecosystem Integrations

- Ambari
 - Support a fully managed NiFi Cluster through Ambari
 - Monitoring, management, upgrades, etc.
- Ranger
 - Ability to delegate authorization decisions to Ranger
 - Manage authorization controls through Ranger
- Atlas
 - Track lineage from the source to destination
 - Apply tags to data as its acquired

The Future – Apache NiFi

- HA Control Plane
 - Zero Master cluster, Web UI accessible from any node
 - Auto-Election of “Cluster Coordinator” and “Primary Node” through ZooKeeper
- HA Data Plane
 - Ability to replicate data across nodes in a cluster
- Multi-Tenancy
 - Restrict Access to portions of a flow
 - Allow people/groups within an organization to only access their portions of the flow
- Extension Registry
 - Create a central repository of NARs and Templates
 - Move most NARs out of Apache NiFi distribution, ship with a minimal set

The Future – Apache NiFi

- Variable Registry
 - Define environment specific variables through the UI, reference through EL
 - Make templates more portable across environments/instances
- Redesign of User Interface
 - Modernize look & feel, improve usability, support multi-tenancy
- Continued Development of Integration Points
 - New processors added continuously!
- MiNiFi
 - Complimentary data collection agent to NiFi's current approach
 - Small, lightweight, centrally managed agent that integrates with NiFi for follow-on dataflow management

Thank you